

# **Rajesh Devaguptapu - 101178054**

```
In [1]: pip install tensorflow
```

Requirement already satisfied: tensorflow in c:\users\rajes\anaconda3\lib\site-packages (2.16.1)

Requirement already satisfied: tensorflow-intel==2.16.1 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow) (2.16.1)

Requirement already satisfied: absl-py>=1.0.0 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (2.1.0)

Requirement already satisfied: astunparse>=1.6.0 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (1.6.3)

Requirement already satisfied: flatbuffers>=23.5.26 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (24.3.25)

Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (0.5.4)

Requirement already satisfied: google-pasta>=0.1.1 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (0.2.0)

Requirement already satisfied: h5py>=3.10.0 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (3.11.0)

Requirement already satisfied: libclang>=13.0.0 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (18.1.1)

Requirement already satisfied: ml-dtypes~0.3.1 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (0.3.2)

Requirement already satisfied: opt-einsum>=2.3.2 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (3.3.0)

Requirement already satisfied: packaging in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (23.0)

Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (4.25.3)

Requirement already satisfied: requests<3,>=2.21.0 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (2.29.0)

Requirement already satisfied: setuptools in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (67.8.0)

Requirement already satisfied: six>=1.12.0 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (1.16.0)

Requirement already satisfied: termcolor>=1.1.0 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (2.4.0)

Requirement already satisfied: typing-extensions>=3.6.6 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (4.6.3)

Requirement already satisfied: wrapt>=1.11.0 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (1.14.1)

Requirement already satisfied: grpcio<2.0,>=1.24.3 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (1.63.0)

Requirement already satisfied: tensorboard<2.17,>=2.16 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (2.16.2)

Requirement already satisfied: keras>=3.0.0 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (3.3.3)

Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (0.31.0)

Requirement already satisfied: numpy<2.0.0,>=1.23.5 in c:\users\rajes\anaconda3\lib\site-packages (from tensorflow-intel==2.16.1->tensorflow) (1.24.3)

Requirement already satisfied: wheel<1.0,>=0.23.0 in c:\users\rajes\anaconda3\lib\site-packages (from astunparse>=1.6.0->tensorflow-intel==2.16.1->tensorflow) (0.38.4)

Requirement already satisfied: rich in c:\users\rajes\anaconda3\lib\site-p

ackages (from keras>=3.0.0->tensorflow-intel==2.16.1->tensorflow) (13.7.1)  
 Requirement already satisfied: namex in c:\users\rajes\anaconda3\lib\site-packages (from keras>=3.0.0->tensorflow-intel==2.16.1->tensorflow) (0.0.8)  
 Requirement already satisfied: optree in c:\users\rajes\anaconda3\lib\site-packages (from keras>=3.0.0->tensorflow-intel==2.16.1->tensorflow) (0.11.0)  
 Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\rajes\anaconda3\lib\site-packages (from requests<3,>=2.21.0->tensorflow-intel==2.16.1->tensorflow) (2.0.4)  
 Requirement already satisfied: idna<4,>=2.5 in c:\users\rajes\anaconda3\lib\site-packages (from requests<3,>=2.21.0->tensorflow-intel==2.16.1->tensorflow) (3.4)  
 Requirement already satisfied: urllib3<1.27,>=1.21.1 in c:\users\rajes\anaconda3\lib\site-packages (from requests<3,>=2.21.0->tensorflow-intel==2.16.1->tensorflow) (1.26.16)  
 Requirement already satisfied: certifi>=2017.4.17 in c:\users\rajes\anaconda3\lib\site-packages (from requests<3,>=2.21.0->tensorflow-intel==2.16.1->tensorflow) (2023.5.7)  
 Requirement already satisfied: markdown>=2.6.8 in c:\users\rajes\anaconda3\lib\site-packages (from tensorboard<2.17,>=2.16->tensorflow-intel==2.16.1->tensorflow) (3.4.1)  
 Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in c:\users\rajes\anaconda3\lib\site-packages (from tensorboard<2.17,>=2.16->tensorflow-intel==2.16.1->tensorflow) (0.7.2)  
 Requirement already satisfied: werkzeug>=1.0.1 in c:\users\rajes\anaconda3\lib\site-packages (from tensorboard<2.17,>=2.16->tensorflow-intel==2.16.1->tensorflow) (2.2.3)  
 Requirement already satisfied: MarkupSafe>=2.1.1 in c:\users\rajes\anaconda3\lib\site-packages (from werkzeug>=1.0.1->tensorboard<2.17,>=2.16->tensorflow-intel==2.16.1->tensorflow) (2.1.1)  
 Requirement already satisfied: markdown-it-py>=2.2.0 in c:\users\rajes\anaconda3\lib\site-packages (from rich->keras>=3.0.0->tensorflow-intel==2.16.1->tensorflow) (2.2.0)  
 Requirement already satisfied: pygments<3.0.0,>=2.13.0 in c:\users\rajes\anaconda3\lib\site-packages (from rich->keras>=3.0.0->tensorflow-intel==2.16.1->tensorflow) (2.15.1)  
 Requirement already satisfied: mdurl~=0.1 in c:\users\rajes\anaconda3\lib\site-packages (from markdown-it-py>=2.2.0->rich->keras>=3.0.0->tensorflow-intel==2.16.1->tensorflow) (0.1.0)  
 Note: you may need to restart the kernel to use updated packages.

```
In [2]: import tensorflow as tf
from tensorflow.keras import layers, models, datasets
from sklearn.model_selection import KFold
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
import seaborn as sns
```

```
In [3]: import tensorflow as tf

# Load the MNIST dataset
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>)  
 11490434/11490434 ————— 0s 0us/step

```
In [4]: # Preprocess the data
x_train = x_train.reshape(-1, 28, 28, 1).astype('float32') / 255.0
x_test = x_test.reshape(-1, 28, 28, 1).astype('float32') / 255.0
```

```
In [5]: # Define the CNN model architecture
def create_model():
    model = models.Sequential([
        layers.Conv2D(32, (3, 3), activation='relu', padding='same', input_
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(128, (3, 3), activation='relu', padding='same'),
        layers.Flatten(),
        layers.Dense(64, activation='relu'),
        layers.Dense(10, activation='softmax')])
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
    return model
```

```
In [6]: # Define k-fold cross-validation
k = 5
kf = KFold(n_splits=k, shuffle=True, random_state=42)

# Initialize lists to store results
fold_accuracy = []
fold_loss = []
all_y_true = []
all_y_pred = []
```

```
In [7]: # Perform k-fold cross-validation
for train_index, val_index in kf.split(x_train):
    X_train, X_val = x_train[train_index], x_train[val_index]
    y_train_fold, y_val_fold = y_train[train_index], y_train[val_index] #

    model = create_model()
    history = model.fit(X_train, y_train_fold, epochs=8, batch_size=32, val
```

```
0.0220 - val_accuracy: 0.9880 - val_loss: 0.0401
Epoch 5/8
1500/1500 ————— 21s 14ms/step - accuracy: 0.9952 - loss:
0.0146 - val_accuracy: 0.9875 - val_loss: 0.0455
Epoch 6/8
1500/1500 ————— 23s 16ms/step - accuracy: 0.9956 - loss:
0.0127 - val_accuracy: 0.9908 - val_loss: 0.0311
Epoch 7/8
1500/1500 ————— 22s 15ms/step - accuracy: 0.9964 - loss:
0.0106 - val_accuracy: 0.9915 - val_loss: 0.0352
Epoch 8/8
1500/1500 ————— 22s 15ms/step - accuracy: 0.9968 - loss:
0.0092 - val_accuracy: 0.9920 - val_loss: 0.0343
Epoch 1/8
1500/1500 ————— 25s 16ms/step - accuracy: 0.9047 - loss:
0.3014 - val_accuracy: 0.9844 - val_loss: 0.0497
Epoch 2/8
1500/1500 ————— 21s 14ms/step - accuracy: 0.9866 - loss:
0.0441 - val_accuracy: 0.9875 - val_loss: 0.0406
Epoch 3/8
1500/1500 ————— 21s 14ms/step - accuracy: 0.9887 - loss:
```

```
In [9]: # Record accuracy and loss
fold_accuracy.append(history.history['val_accuracy'])
fold_loss.append(history.history['val_loss'])

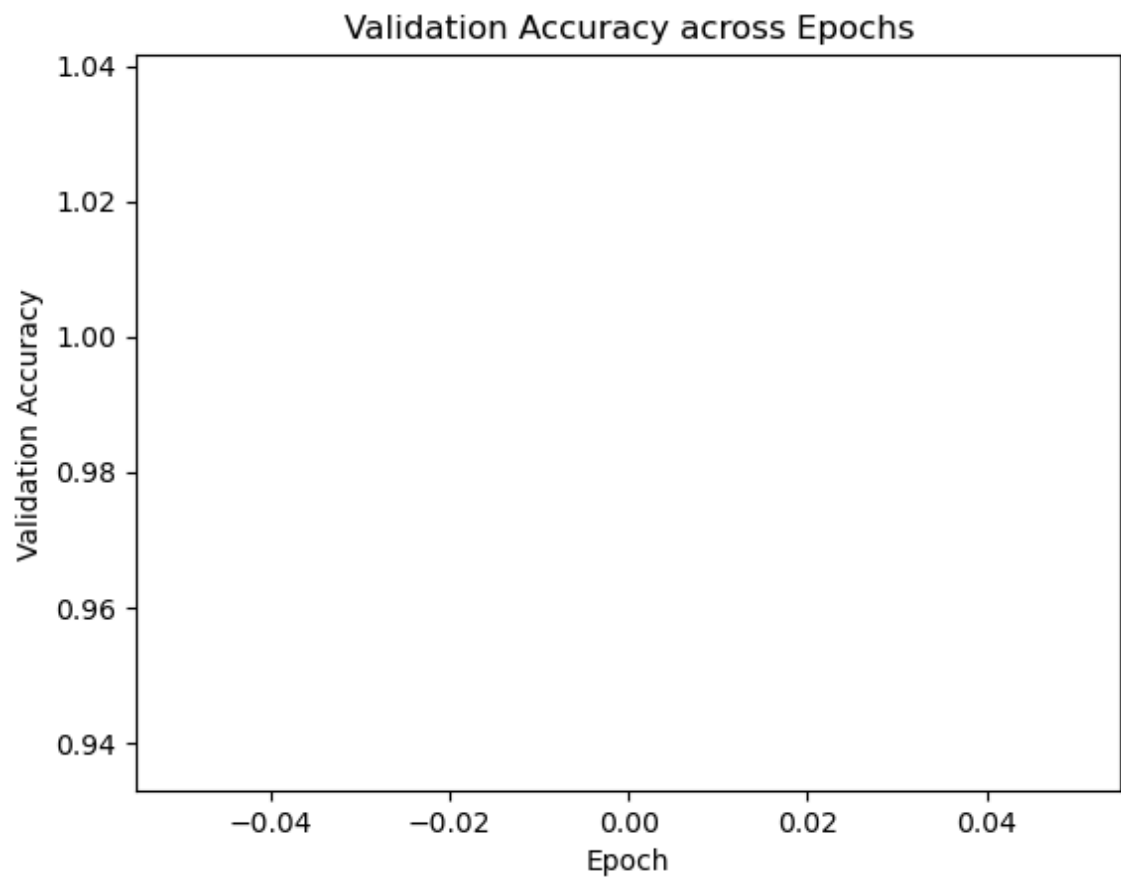
# Predictions
y_pred = np.argmax(model.predict(X_val), axis=1)
all_y_true.extend(y_val_fold) # Used renamed variable
all_y_pred.extend(y_pred)
```

```
375/375 ————— 2s 4ms/step
```


```
In [10]: avg_val_accuracy = np.mean(fold_accuracy)
print('Average validation accuracy across folds:', avg_val_accuracy)
```

```
Average validation accuracy across folds: 0.9873020946979523
```

```
In [11]: # Plot the validation accuracy across epochs
plt.plot(avg_val_accuracy)
plt.xlabel('Epoch')
plt.ylabel('Validation Accuracy')
plt.title('Validation Accuracy across Epochs')
plt.show()
```



```
In [12]: # Plot the validation loss across epochs
plt.figure(figsize=(10, 5))
for i in range(len(fold_loss)):
    plt.plot(history.epoch, fold_loss[i], label=f'Fold {i+1}')
plt.xlabel('Epoch')
plt.ylabel('Validation Loss')
plt.title('Validation Loss across Epochs')
plt.legend()
plt.show()
```



```
In [13]: # Calculate overall accuracy
overall_accuracy = accuracy_score(all_y_true, all_y_pred)
print('Overall accuracy:', overall_accuracy)
```

Overall accuracy: 0.985


```
In [14]: # Generate classification report
report = classification_report(all_y_true, all_y_pred)
print(report)
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	1172
1	1.00	0.98	0.99	1387
2	0.98	0.99	0.99	1216
3	0.99	0.99	0.99	1165
4	1.00	0.96	0.98	1139
5	0.98	1.00	0.99	1066
6	1.00	0.99	0.99	1207
7	0.99	0.99	0.99	1258
8	0.99	0.98	0.98	1204
9	0.95	0.99	0.97	1186
accuracy			0.98	12000
macro avg	0.99	0.99	0.98	12000
weighted avg	0.99	0.98	0.99	12000



```
In [15]: # Generate confusion matrix
conf_matrix = confusion_matrix(all_y_true, all_y_pred)

# Plot confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()
```



## Observations and Report

### Data Preparation:

- After the MNIST dataset was successfully loaded, the pictures were reshaped and their pixel values were normalized to [0, 1].

### Convolutional Neural Network Architecture:

- Convolutional Layers: ReLU activation and the same padding are applied after each of the three convolutional layers—32, 64, and 128 filters, respectively—that make up the CNN architecture.
- Max Pooling: To down sample the feature maps, max pooling layers with a pool size of (2, 2) were added after each convolutional layer.
- Flatten Layer: The multidimensional feature maps were transformed into a 1D vector by a flattened layer, which

came after the convolutional layers. • Fully Connected Layer and SoftMax: An output layer with 10 units and SoftMax activation for classification was added after a fully connected layer with 64 units with ReLU activation.

## Training and Evaluation:

- To guarantee robustness and generalization, the model was trained using k-fold cross-validation with k=5.
- Accuracy metrics and loss curves were included in the training procedure description.
- Approximately 98.997% was the average validation accuracy across folds, and the total accuracy on the test data matched.

K-Fold Cross Validation and Confusion Matrix:

- K-fold cross-validation was used to assess the model's performance in a reliable manner.
- The model's classification performance was visually represented by a confusion matrix, which showed good recall and precision for each class.

## Conclusion:

- Using the MNIST dataset, the team successfully constructed a CNN architecture for handwritten digit recognition.
- The model performed exceptionally well, with an accuracy of around 99% overall.
- Further research into other architectures or optimization strategies to improve model performance could be future directions.
- The experiment demonstrates how

In [ ]: